

A Hybrid ARIMA-LSTM Framework for Remaining Useful Life Prediction in Industrial IoT Sensor Data: A Cross-Regime Study on NASA C-MAPSS

Pradeep Somasundaram
Independent Researcher
Dallas, TX, USA
aadhi1501@gmail.com

Abstract—Remaining useful life (RUL) prediction is central to prognostics and health management in industrial IoT systems, where sensor streams from rotating machinery combine slow degradation trends with sharper, nonlinear fault dynamics. One way to help a sequence model separate these two components is to fit a classical time-series model to each sensor channel and give the network both the raw signal and the fitted trend, an idea with some support in prior work on other kinds of equipment. We test it on NASA’s C-MAPSS turbofan benchmark, and we test it twice: once with the fixed ARIMA(1,1,0) order common in prior work, and once with a per-channel, AIC-selected order, comparing both hybrid variants against a plain LSTM and a Random Forest baseline on all four C-MAPSS sub-datasets, each run under six random seeds. Under a fixed order, the picture is negative and, on its own, not fully stable: mean RMSE is at or above the plain LSTM’s on every sub-dataset, significantly so on FD004 ($p = 0.004$), and an earlier three-seed version of the same comparison had instead flagged FD001, not FD004, as significant, a reversal we report directly rather than fold into a cleaner-looking narrative. Under a properly selected order, the result changes: the selected-order hybrid significantly and consistently outperforms the fixed-order hybrid on two of four sub-datasets (FD002, FD004; $p = 0.031$, 6 of 6 seeds each), and on FD002 it also significantly outperforms the plain LSTM baseline itself (13.77 vs. 14.59 cycles RMSE, $p = 0.031$, 6 of 6 seeds), the one result in this paper we consider a genuine, seed-robust win for the trend-augmented architecture. The honest summary is not that trend augmentation fails, but that it fails under the field’s default fixed-order convention and can succeed, robustly, once that convention is checked rather than assumed, a distinction we could only draw because we insisted on six-seed replication for every claim, including the ones that first looked negative.

Index Terms—predictive maintenance, remaining useful life, LSTM, ARIMA, prognostics and health management, industrial IoT, big data, multi-condition degradation

I. INTRODUCTION

Unplanned downtime in industrial and manufacturing settings carries substantial operational and safety cost, motivating a large body of work on prognostics and health management (PHM), the discipline of estimating a system’s remaining useful life (RUL) from streaming sensor data so that maintenance can be scheduled proactively rather than reactively [1]. Industrial IoT deployments now produce continuous, high-volume,

multivariate sensor telemetry (vibration, temperature, pressure, rotational speed) from equipment such as turbines, pumps, and CNC machinery, placing RUL estimation squarely in the big-data regime: large sensor volume, high update velocity, and heterogeneous multivariate structure. A recent survey by Khan et al. [10] catalogs the resulting breadth of deep-learning-based PdM methods (CNNs, autoencoders, GANs, transformers, and foundation models) across supervised, transfer, and federated learning paradigms, which gives some sense of how active, and how fragmented, this area has become.

Two modeling traditions address the RUL estimation problem in different ways. Classical time-series models, most notably the Autoregressive Integrated Moving Average (ARIMA) family, are well suited to capturing the smooth trend and autocorrelation structure of a single degrading signal, but assume linearity and handle multivariate sensor interaction poorly. Deep sequence models, particularly Long Short-Term Memory (LSTM) networks [3], learn nonlinear temporal dependencies directly from raw multivariate sensor windows and have become a standard approach for RUL estimation, but a randomly initialized LSTM has to learn both the trend and the nonlinear fault dynamics from data alone, which can be inefficient, particularly for short degradation histories or noisy multi-condition operation. It is a natural idea to try to give the LSTM a head start on the trend part.

This paper tests that idea directly: for each sensor channel we fit a low-order per-unit ARIMA model and use its fitted, in-sample one-step values as an explicit trend feature, concatenated with the raw sensor readings before the resulting multivariate sequence goes into an LSTM regressor. The hope is that this relieves the LSTM of re-deriving the smooth trend component from scratch, letting it spend its representational capacity on the residual nonlinear dynamics and cross-sensor interactions instead.

We evaluate this across all four C-MAPSS sub-datasets, spanning three qualitatively distinct operating regimes: single condition, single fault mode (FD001); single condition, two fault modes (FD003); multi-condition, single fault mode (FD002); and multi-condition, two fault modes (FD004). More

importantly, after an early result made us suspicious of our own conclusions, we do not stop at a single training run, or even at a single small batch of them. We repeat every comparison across six random seeds, and separately check whether a fixed ARIMA order was itself the problem by re-running a per-channel order-selection variant. Both checks matter, and we report what changed each time we added more evidence, including cases where the added evidence overturned a finding we had already started to explain.

Our contributions are: (1) a lightweight, per-sensor ARIMA trend-augmentation scheme for LSTM-based RUL regression that adds no additional labeled-data requirements; (2) an evaluation across all four C-MAPSS sub-datasets, using regime-aware normalization [4] for the multi-condition subsets, against plain-LSTM and Random Forest baselines, each averaged over six seeds; (3) a second robustness check using per-channel ARIMA order selection in place of a fixed order, which turns out to be just as seed-sensitive as the main comparison; (4) a properly seed-averaged window-length ablation; and (5) a direct, documented account of how this paper’s own headline finding changed as we added more seeds, offered as a concrete illustration of why single-seed and even small-multi-seed comparisons on this benchmark deserve some skepticism.

II. RELATED WORK

RUL estimation on C-MAPSS. The NASA C-MAPSS turbofan degradation dataset [1] is the standard benchmark for data-driven RUL estimation. Ramasso and Saxena [2] laid out early benchmarking protocols across similarity-based and regression approaches. More recent work has moved toward deep sequence and hybrid architectures: Muthukumar and Philip [5] combine convolutional feature extraction with LSTM temporal modeling, and Chowdhury et al. [6] propose a residual-corrected bidirectional LSTM that separates a coarse trend prediction from a learned residual correction, a decomposition philosophy related to, but architecturally distinct from, the ARIMA-based decomposition used here.

Multi-condition RUL estimation. FD002 and FD004 introduce six discrete operating regimes and, in FD004’s case, two concurrent fault modes, both of which increase estimation difficulty relative to FD001’s single-condition setting. Su and Fang [7] propose a multi-head attention fusion network that decomposes degradation into trend, operational-state embedding, and residual-noise components under discrete operating conditions, a decomposition philosophy that motivates our own trend/residual split, though implemented with ARIMA rather than a learned embedding.

Hybrid ARIMA-deep learning models. Outside of turbofan RUL specifically, ARIMA-LSTM hybrids have been applied to remaining-useful-life estimation of water-hydraulic high-speed on/off valves [8], where ARIMA models the linear degradation trend and an auxiliary classifier detects health-state transitions, and to battery capacity and life estimation [9], where ARIMA-smoothed inputs feed an attention-based CNN-LSTM. Both report the hybrid approach helping on their

respective single-regime datasets. Neither paper states how many random seeds its headline numbers are based on. Given that our own headline finding changed twice as we added seeds and re-checked the ARIMA order (Sections V and VI), we do not think this detail is minor, and we raise it here as an open question about the hybrid-forecasting literature more broadly rather than a specific claim about either paper.

Reproducibility and seed sensitivity in machine learning.

The instability we document in Section V is a specific instance of a broader problem that other subfields of machine learning have already named. Henderson et al. [12] showed that headline deep reinforcement learning results routinely fail to replicate across random seeds and codebase details, and proposed reporting practices, including seed counts and variance, that much of the field still does not follow. Lipton and Steinhardt [13] catalog related scholarship failures, including a failure to distinguish which part of an empirical gain a given change is actually responsible for, which is close to the mistake our own three-seed FD001 result nearly encoded into this paper’s discussion. Bouthillier et al. [14] make the point most directly relevant to our own experience: they show experimentally that a result can be reproducible, in the sense that rerunning the exact same configuration gives the same answer, while still not being robust to small, legitimate changes in experimental setup, which is exactly the relationship between our fixed-seed and multi-seed FD001 results. To our knowledge, this specific caution has not previously been raised inside the C-MAPSS RUL literature, where fixed single-seed comparisons are the norm rather than the exception. A concrete, recent example: Goel et al. [15], published on this exact benchmark within weeks of our own experiments, state plainly that “all experiments use random seed 42 for reproducibility.” We do not single this out to criticize that paper specifically, single-seed evaluation is simply the field’s default, not an unusual choice, but the phrasing is a precise illustration of the conflation Bouthillier et al. [14] warn about: a fixed seed makes a result exactly reproducible, in the narrow sense that rerunning it gives the same number, without making it representative of what the method does on average, which is the property a reader actually cares about. We see our contribution here as importing an already-known methodological concern into a subfield that has not yet widely adopted it, not as identifying the concern for the first time.

Our approach differs from prior turbofan RUL work in that the ARIMA component is used purely as a per-channel feature-engineering step, producing an explicit trend covariate per sensor, rather than as a separate forecaster whose output is fused post-hoc with a deep model’s prediction, and in that we evaluate it across all four regimes, across seeds, and across two different ARIMA order-selection strategies rather than on a single convenient run.

III. METHODOLOGY

A. Datasets

We use all four NASA C-MAPSS sub-datasets [1]: **FD001** (100 train / 100 test engines, 1 operating condition, 1 fault

mode: HPC degradation), **FD002** (260 train / 259 test engines, 6 operating conditions, 1 fault mode), **FD003** (100 train / 100 test engines, 1 operating condition, 2 fault modes: HPC and fan degradation), and **FD004** (249 train / 248 test engines, 6 operating conditions, 2 fault modes). Each engine unit reports 3 operational settings and 21 sensor channels per cycle; training units run to failure while test units are truncated at an arbitrary cycle, and the task is to predict the true remaining cycles at that truncation point.

B. Preprocessing

Following standard practice [2], we cap the training RUL label at 125 cycles, reflecting the empirical observation that engine degradation is roughly constant before this point and roughly linear afterward. We discard sensor channels with near-zero variance in the training set: 6 channels for FD001 and FD003 (15 informative channels retained for each), and none for FD002/FD004 (all 21 channels retain variance, since operating-condition changes alone introduce variance in otherwise-flat channels).

Regime-aware normalization. FD002 and FD004 combine six discrete operating regimes, which shift the baseline value of most sensors independently of degradation state and would otherwise confound a global normalization scheme. Following Li, Ding, and Sun [4], we cluster the three-dimensional operational-setting vectors into K regimes with K -means fit on the training set ($K=6$ for FD002/FD004; $K=1$ for FD001/FD003, which reduces to ordinary global normalization, so the same code path is used for all four sub-datasets), assign each test-set row to its nearest regime centroid, and z -score normalize each sensor channel using regime-specific training-set statistics, clipping to ± 6 standard deviations to control outliers in sparsely populated regimes.

C. ARIMA Trend-Feature Extraction

For each engine unit and each retained sensor channel, we independently fit an ARIMA model to that channel's cycle-indexed, regime-normalized time series. Writing y_t for a given channel's reading at cycle t , an ARIMA(1,1,0) model, our main-comparison default, treats the first difference $\Delta y_t = y_t - y_{t-1}$ as an AR(1) process,

$$\Delta y_t = \phi \Delta y_{t-1} + \varepsilon_t, \quad \varepsilon_t \sim \mathcal{N}(0, \sigma^2), \quad (1)$$

equivalently $y_t = y_{t-1} + \phi(y_{t-1} - y_{t-2}) + \varepsilon_t$. We fit ϕ and σ^2 by maximum likelihood independently per (unit, sensor) pair and take the in-sample one-step-ahead fitted value $\hat{y}_t$, aligned index-for-index with the original reading, as the trend feature. Where a fit fails to converge for a short series, we fall back to a 5-cycle rolling mean.

Order selection variant. Section V-C reports a second version of this feature where, instead of fixing the order at (1, 1, 0), we fit four candidate models per (unit, sensor) pair, (1, 1, 0), (0, 1, 1), (1, 1, 1), and (2, 1, 0), and keep the fitted values from whichever has the lowest Akaike Information Criterion (AIC). Everything downstream, windowing, the LSTM architecture, training procedure, is identical between the two variants; only the trend feature's construction differs.

D. Hybrid Model Architecture

We construct sliding windows of length W cycles ($W=30$ in our main experiments; see Section V-D for a window-length ablation) over the concatenation of raw sensor features and their ARIMA trend features. Each window is passed through a 2-layer LSTM (hidden size 64) whose cell updates follow the standard formulation [3]

$$\begin{aligned} f_t &= \sigma(W_f[h_{t-1}, x_t] + b_f), & i_t &= \sigma(W_i[h_{t-1}, x_t] + b_i) \\ \tilde{c}_t &= \tanh(W_c[h_{t-1}, x_t] + b_c), & o_t &= \sigma(W_o[h_{t-1}, x_t] + b_o) \\ c_t &= f_t \odot c_{t-1} + i_t \odot \tilde{c}_t, & h_t &= o_t \odot \tanh(c_t) \end{aligned} \quad (2)$$

followed by a feed-forward regression head ($64 \rightarrow 32 \rightarrow 1$) predicting capped RUL from the final hidden state h_W . Training minimizes mean squared error,

$$\mathcal{L} = \frac{1}{N} \sum_{i=1}^N \left(\text{RUL}_i - \widehat{\text{RUL}}_i \right)^2, \quad (3)$$

with Adam, using an engine-level held-out validation split (15% of training units, never seen during training) for early stopping.

E. Baselines

We compare against (i) a *Random Forest* regressor (300 trees, max depth 10) trained on per-cycle flat feature snapshots, every training cycle rather than just each engine's terminal cycle, since training on terminal cycles alone would teach a degenerate near-zero-RUL target given that training engines run to failure, as a non-sequential baseline; and (ii) a *plain LSTM* with identical architecture and training procedure to the hybrid model but using only the raw sensor features, isolating the effect of the trend features.

F. Evaluation Protocol

Following the official C-MAPSS test protocol, each test engine's RUL is predicted once, from its final available window, left-padded by repetition for engines with fewer than W recorded cycles, and compared against the ground-truth RUL supplied with the test split. We report RMSE, MAE, and the asymmetric PHM scoring function of [1]: writing $d_i = \widehat{\text{RUL}}_i - \text{RUL}_i$,

$$s_i = \begin{cases} \exp(-d_i/13) - 1, & d_i < 0 \text{ (early/conservative)} \\ \exp(d_i/10) - 1, & d_i \geq 0 \text{ (late/over-optimistic)} \end{cases} \quad (4)$$

with total score $\sum_i s_i$, penalizing late predictions more heavily than early ones, since failing to service a machine in time is costlier than servicing it early.

IV. EXPERIMENTAL SETUP

All experiments use PyTorch (CPU) and statsmodels for ARIMA fitting, scikit-learn for K -means regime clustering and the Random Forest baseline, and SciPy for the significance tests in Section V-B. Table I reports the full LSTM hyperparameter configuration, and Table II reports dataset scale.

The main cross-dataset comparison in Table III is repeated under six seeds (42, 123, 2024, 7, 99, 314); we report

TABLE I
LSTM HYPERPARAMETER CONFIGURATION (IDENTICAL FOR PLAIN AND HYBRID MODELS; ONLY INPUT WIDTH DIFFERS, 15–21 RAW CHANNELS VS. DOUBLE THAT WITH ARIMA TREND CHANNELS APPENDED).

Hyperparameter	Value
LSTM layers	2
Hidden size	64
Recurrent dropout	0.2
Regression head	64 $\rightarrow$ 32 $\rightarrow$ 1 (ReLU)
Optimizer	Adam, lr = 10^{-3}
Loss	MSE on capped RUL
Batch size	128
Max epochs	60
Early-stopping patience	8 epochs (main runs)
Window length W	30 cycles (main runs)
RUL cap	125 cycles
ARIMA order (main runs)	(1, 1, 0), fixed per channel
ARIMA order (Sec. V-C)	AIC-selected from 4 candidates
Validation split	15% of training engines
Seeds (main comparison)	42, 123, 2024, 7, 99, 314
Seeds (ablation)	42, 123, 2024

TABLE II
DATASET SCALE AND APPROXIMATE SINGLE-SEED TRAINING RUNTIME.

Dataset	Regimes	Train units	Train windows	Runtime
FD001	1	85	15,104	85s
FD002	6	221	39,513	210s
FD003	1	85	18,600	130s
FD004	6	212	46,462	265s

mean and standard deviation over these six runs, and for the significance test in Section V-B we average each engine’s error across seeds before testing. We did not settle on six seeds in advance as a planned sample size. We started with three, found a significant result on FD001, wrote up an explanation for it, and then, out of caution, added three more seeds before treating that explanation as final. The three additional seeds erased the FD001 result and produced a different significant result on FD004 instead. We report both the three-seed and six-seed pictures in Section V-A rather than only the final one, because the change between them is, in our view, more informative than either snapshot alone.

V. RESULTS

A. Cross-Dataset Results

Table III reports RMSE, MAE, and PHM score, mean $\pm$ standard deviation over six seeds, for all three models across all four datasets, and Fig. 1 shows the same RMSE comparison with error bars.

Sequence modeling remains a clear win over the non-sequential Random Forest baseline everywhere, a 12–27% RMSE reduction depending on the dataset. The comparison between plain and hybrid LSTM again favors the plain model on average, on all four datasets, though the margins are smaller with six seeds than they first appeared with three. FD003 is close to a tie (14.87 vs. 14.72); the other three datasets show the hybrid model’s mean RMSE between 2% and 5% higher than the plain model’s.




Fig. 1. RMSE comparison across all four C-MAPSS sub-datasets, mean $\pm$ one standard deviation over six seeds. The plain LSTM’s mean RMSE is at or below the hybrid model’s on every dataset; note FD003’s wide error bars, driven mostly by the plain model’s own across-seed variance there, discussed in Section VI.

What changed between three seeds and six. We want to walk through this in some detail because we think it is the most useful thing in the paper. With the first three seeds (42, 123, 2024), the hybrid model’s mean RMSE was higher than the plain model’s on FD001, FD002, and FD004, and a paired Wilcoxon test on per-engine error reached significance on FD001 ($p = 0.007$) and nowhere else. We wrote a discussion section around that: a redundancy argument for why FD001’s clean single-condition, single-fault setting would show the clearest effect. Adding three more seeds (7, 99, 314) changed the specific answer. FD001’s significance disappeared ($p = 0.157$ at six seeds); FD004, which had not been significant at three seeds ($p = 0.159$), became significant at six ($p = 0.004$, Section V-B). The direction of the overall finding, hybrid mean RMSE at or above plain on every dataset, held up both times. Which dataset you could defensibly call out as a significant, specific effect did not.

For outside context, Sharma [11] reports RMSE of 16.22, 19.29, 16.84, and 19.98 on FD001 through FD004 with an attention-based bidirectional LSTM. Our plain-LSTM baseline, with no attention mechanism and no trend feature, is lower on all four (14.43, 14.59, 14.72, 14.92). We would not lean too hard on this comparison, since the two setups may not share the same RUL cap, window length, or train/validation protocol, but it is at least a sign that our baseline is not a weak strawman.

B. Statistical Significance

Table III’s differences are averages over 100–259 test engines and six seeds; a directional gap in the mean does not by itself tell us whether the two models make systematically different errors on individual engines. We run a paired Wilcoxon signed-rank test on per-engine absolute error, hybrid vs. plain LSTM, using each engine’s error averaged across the six seeds (Table IV).

At six seeds, FD004 is the only dataset that reaches significance at $\alpha=0.05$: the hybrid model is worse on 144 of 248 engines, with a median per-engine increase in absolute error of 0.41 cycles. FD001 and FD002 show the same direction

TABLE III

RUL PREDICTION PERFORMANCE ACROSS NASA C-MAPSS SUB-DATASETS, MEAN $\pm$ STD OVER 6 SEEDS (42, 123, 2024, 7, 99, 314). TEST ENGINES: FD001/FD003 $n=100$, FD002 $n=259$, FD004 $n=248$.

Dataset	Model	RMSE	MAE	PHM Score
FD001 (1 condition, 1 fault)	Random Forest	17.34 ± 0.28	12.23 ± 0.18	936.9 ± 97.3
	Plain LSTM	14.43 ± 0.64	10.43 ± 0.66	505.9 ± 168.6
	Hybrid ARIMA-LSTM	14.92 ± 1.65	11.04 ± 1.19	478.8 ± 210.8
FD002 (6 conditions, 1 fault)	Random Forest	17.97 ± 0.12	13.39 ± 0.08	3383.6 ± 190.3
	Plain LSTM	14.59 ± 0.49	10.17 ± 0.34	1204.0 ± 111.7
	Hybrid ARIMA-LSTM	14.93 ± 0.60	10.53 ± 0.28	1360.3 ± 280.2
FD003 (1 condition, 2 faults)	Random Forest	20.06 ± 0.36	14.78 ± 0.22	1748.2 ± 148.5
	Plain LSTM	14.72 ± 2.06	10.27 ± 1.44	626.9 ± 414.2
	Hybrid ARIMA-LSTM	14.87 ± 0.58	10.38 ± 0.47	571.1 ± 270.9
FD004 (6 conditions, 2 faults)	Random Forest	20.09 ± 0.09	14.67 ± 0.08	3770.5 ± 123.7
	Plain LSTM	14.92 ± 0.38	10.59 ± 0.32	1112.7 ± 130.9
	Hybrid ARIMA-LSTM	15.71 ± 0.54	11.11 ± 0.50	1342.5 ± 299.7

TABLE IV

PAIRED WILCOXON SIGNED-RANK TEST ON PER-ENGINE ABSOLUTE ERROR, AVERAGED OVER 6 SEEDS (HYBRID VS. PLAIN LSTM). THE FD001 ROW FOR 3 SEEDS IS SHOWN FOR COMPARISON; SEE SECTION V-A.

Dataset	Seeds	p -value	Hybrid worse on	Median $\Delta \text{err} $
FD001	3	0.007	62/100	+0.91
FD001	6	0.157	52/100	-0.09
FD002	6	0.142	134/259	-0.07
FD003	6	0.880	50/100	+0.01
FD004	6	0.004	144/248	-0.41




Fig. 2. Predicted vs. true RUL for the hybrid model, pooling all six seeds, across all four test sets (dashed line = perfect prediction).

(hybrid nominally worse) without reaching significance, and FD003 is close to flat.

Because Table IV reports four tests, one per dataset, rather than one pre-registered test, we also check whether FD004’s result survives correction for multiple comparisons. Applying the Holm-Bonferroni step-down procedure across the four six-seed p -values (0.004, 0.142, 0.157, 0.880), FD004’s smallest p -value is compared against $\alpha/4 = 0.0125$ rather than $\alpha = 0.05$; at $p = 0.004$, it clears this stricter threshold, and no other dataset comes close to its own step-down threshold. FD004 is therefore the one result in this paper we are willing to call significant under a family-wise error rate correction, not only under an uncorrected per-dataset test. Fig. 3 shows the underlying seed-averaged error distributions for all four datasets side by side; as with the aggregate table, the FD003 plain-model distribution is visibly wider than the others, which we return to in Section VI.




Fig. 3. Per-engine absolute error, averaged over six seeds, plain vs. hybrid LSTM, across all four datasets.

C. ARIMA Order Selection

The main comparison fixes the ARIMA order at $(1, 1, 0)$ for every sensor and every dataset. This is a modeling choice, not a finding, and it is worth asking whether the hybrid model’s results are an artifact of that specific choice rather than of trend augmentation itself. We re-ran the ARIMA feature extraction with per-channel AIC-based order selection among four candidates (Section III-C) and re-trained both LSTM variants on the resulting features, at all six seeds across all four datasets, the same replication we used for the main comparison, after an early two-seed version of this check (reported in an earlier draft of this paper) suggested the answer might itself be seed-dependent and was too thin to trust on its own.

Order selection overwhelmingly preferred $(0, 1, 1)$ over our fixed default: across all four datasets, $(0, 1, 1)$ was chosen for 75–83% of (unit, sensor, seed) combinations, with $(1, 1, 1)$ a distant second (14–20%) and $(1, 1, 0)$, our original default, chosen for at most 4%. Table V reports the resulting hybrid-model RMSE, mean $\pm$ std over six seeds, alongside the fixed-order hybrid and the plain LSTM.

Unlike the two-seed version we started with, this is not a wash. A paired Wilcoxon test, selected order vs. fixed order, reaches significance on FD002 and FD004 ($p = 0.031$ on both, the smallest attainable p -value at $n = 6$ paired samples, with the selected order better on 6 of 6 seeds each time) and does not on FD001 ($p = 0.688$, better on 4 of 6) or FD003 ($p = 0.219$, better on 5 of 6). The more direct question,

TABLE V
HYBRID MODEL RMSE WITH AIC-SELECTED VS. FIXED ARIMA ORDER,
MEAN $\pm$ STD OVER 6 SEEDS. PLAIN LSTM SHOWN FOR REFERENCE.

Dataset	Plain	Fixed order	Selected order
FD001	14.43 $\pm$ 0.64	14.92 $\pm$ 1.65	14.38 $\pm$ 0.55
FD002	14.59 $\pm$ 0.49	14.93 $\pm$ 0.60	13.77$\pm$0.23
FD003	14.72 $\pm$ 2.06	14.87 $\pm$ 0.58	14.45 $\pm$ 0.71
FD004	14.92 $\pm$ 0.38	15.71 $\pm$ 0.54	14.90 $\pm$ 0.42

TABLE VI
WINDOW-LENGTH ABLATION ON FD001, MEAN $\pm$ STD OVER 3 SEEDS
(TEST RMSE, CYCLES).

W (cycles)	10	20	30	50
Plain	18.80 $\pm$ 0.47	17.19 $\pm$ 0.11	14.56 $\pm$ 0.75	14.05$\pm$0.44
Hybrid	18.70 $\pm$ 0.67	16.75$\pm$0.34	13.84$\pm$0.90	14.70 $\pm$ 0.42

whether the order-selected hybrid beats the *plain* LSTM rather than just the fixed-order hybrid, gives a sharper answer still: on FD002, it does, significantly ($p = 0.031$, 6 of 6 seeds, 13.77 vs. 14.59), which is the one place in this paper where the hybrid architecture beats the plain baseline with the same statistical confidence we used to declare FD004’s fixed-order disadvantage significant in Section V-B. On FD004, order selection erases the fixed-order disadvantage but does not turn it into an advantage (14.90 vs. 14.92 plain, $p = 1.000$, essentially a tie). FD001 and FD003 stay ties either way.

Read together with Section V-B, the honest summary is not “the hybrid model does not help” but “the hybrid model does not help *under a naively fixed ARIMA order*, and on at least one dataset, FD002, a properly selected order turns that into a real, seed-robust improvement.” We did not expect this when we started the order-selection check; our working assumption going in, informed by the two-seed version, was that it would turn out to be more noise. Six seeds said otherwise on two of four datasets.

D. Window-Length Ablation

Table VI and Fig. 4 report a window-length ablation on FD001 ($W \in \{10, 20, 30, 50\}$), averaged over three seeds (42, 123, 2024), each trained with a shorter early-stopping patience than the main experiments for faster iteration across four window sizes. Unlike the main cross-dataset comparison, this pattern survives three-seed averaging reasonably well: the hybrid model’s advantage over the plain LSTM grows from essentially nothing at $W=10$ (0.11 cycles) to its largest gap at $W=30$ (0.72 cycles), then reverses at $W=50$, where the hybrid model is 0.65 cycles worse. We still would not call this a fully confirmed effect, it rests on only three seeds and one dataset, but it is the one place in this paper where averaging over seeds did not wash the effect out, and we report it with correspondingly less hedging than the rest of our results.

VI. DISCUSSION

The headline finding is about evaluation, and it points to a real fix. We went into this project expecting a yes-or-no




Fig. 4. FD001 test RMSE vs. window length, mean $\pm$ std over three seeds. The hybrid model’s advantage peaks at $W=30$ and reverses at $W=50$; this is the one pattern in the paper that held up under seed-averaging.

answer to whether a trend feature helps an LSTM predict RUL. Under a fixed ARIMA order, the specific answer moved as we checked our own work more carefully: at three seeds, FD001 looked like a clean, significant, explainable disadvantage for the hybrid model; at six seeds, that result was gone and FD004 had taken its place. We take that reversal seriously, and it is a domain-specific instance of the pattern Henderson et al. [12] and Bouthillier et al. [14] document in reinforcement learning: a configuration can be exactly reproducible and still not be representative of what the method does on average. But unlike our earlier, more tentative draft of this section, we do not think the story ends there. Once we powered the ARIMA order-selection check to the same six seeds as the main comparison, it stopped looking like a third source of noise and started looking like an explanation: on FD002, replacing the fixed order with a properly AIC-selected one turns a nominal disadvantage into a significant, 6-of-6-seed advantage over the plain LSTM (Section V-C). On FD004, it erases the fixed-order disadvantage entirely. On FD001 and FD003, it does neither, and the comparison stays a wash.

What we think is actually going on. An ARIMA(1,1,0) model, fit to a 30-cycle window, is a specific and fairly restrictive assumption, a single autoregressive lag on the first difference. Our own order-selection results show that assumption was wrong for this data 75–83% of the time; (0, 1, 1), a moving-average model on the differenced series, fit better almost everywhere. A trend feature built from a poorly-fit model is not neutral, it is actively misleading, effectively noise correlated with the raw signal it is supposed to summarize, and Section V-C’s fixed-vs-selected comparison suggests that noise is exactly what was suppressing the hybrid model’s performance on FD002 and FD004. Once the trend feature is well-fit, on FD002 at least, it appears to supply real information the plain LSTM was not otherwise recovering from a 30-cycle raw window, consistent with the window-length ablation (Section V-D), where the hybrid model’s advantage is concentrated in a middle range of window lengths and disappears once the window is long enough for the LSTM to reconstruct the trend

unaided. We do not have a mechanistic account of why FD002 specifically, and not FD001 or FD003, is where a well-fit trend feature pays off, and we are not going to force one; we flag it as the most interesting open question this paper raises rather than close it with a guess.

FD003’s variance is worth a separate note. The plain LSTM’s across-seed standard deviation on FD003 (2.06 RMSE) is the largest of any model-dataset pair in Table III, larger than any hybrid-model variance we observed anywhere, fixed-order or selected-order. We do not have a confident explanation for it, and we would rather flag it as unexplained than invent one to fill the gap.

What we would tell someone deciding whether to add this kind of feature to their own pipeline. Do not treat a single run, or even three seeds, as sufficient evidence either way, in either direction. If you are going to report that a trend-augmentation feature helps or hurts, budget for the kind of multi-seed and order-selection re-checks we describe here. Our own answer changed twice before it stabilized into something we trust: first a nominal loss, then a statistically framed loss that itself moved between two sets of three seeds, and only at full six-seed power on both the main comparison and the order-selection check did a consistent, actionable picture emerge, that a naively fixed ARIMA order actively hurts this architecture, and that fixing the order, not abandoning the idea, is where the real effect turned out to live.

Concretely, three checks are cheap enough that we think they should be routine before anyone reports a trend-augmented RUL model as an improvement, or a failure, on a C-MAPSS-scale benchmark. First, rerun the comparison under at least five or six seeds rather than one to three; Section V-A shows this can change which dataset, if any, reaches significance, not just how confident you are in the direction. Second, apply a multiple-comparison correction, Holm-Bonferroni or similar, whenever the comparison is repeated across more than one dataset or condition, as we do in Section V-B; a single uncorrected $p < 0.05$ out of four tests is a weaker claim than it looks. Third, treat a fixed ARIMA order as a modeling choice to be checked, not assumed; Section V-C shows that this check, run at full seed power, is not a formality, it is where two of our four datasets’ real result was hiding.

VII. LIMITATIONS AND THREATS TO VALIDITY

The window-length ablation (Section V-D) uses three seeds on one dataset (FD001) and a shorter training budget than the main experiments; it is the most seed-stable result in the paper, but it is still a single-dataset finding we have not attempted to replicate on FD002–FD004, or under the AIC-selected order. Our comparison against Sharma [11] in Section V-A is informal, since we do not know whether that paper’s RUL cap, window length, or data split are close enough to ours to call it a controlled comparison. We tried four ARIMA orders and one LSTM architecture throughout; we did not sweep LSTM hidden size, depth, or an alternative fusion mechanism (e.g., gating or an auxiliary trend-prediction loss instead of feature concatenation), any of which could plausibly interact with

the order-selection effect in Section V-C. We do not have an explanation for FD003’s unusually high plain-model variance across seeds, and we are reporting it as an open observation rather than resolving it. We also do not have a mechanistic account of why order selection produces a significant, seed-robust win specifically on FD002 and not on FD001 or FD003; our redundancy-based explanation for the fixed-order results does not by itself predict this asymmetry, and we flag that gap rather than paper over it. Finally, every experiment here uses NASA’s simulated turbofan data, and whether any of these patterns extend to real industrial IoT sensor streams with different noise and missingness characteristics is a question we are not in a position to answer.

VIII. CONCLUSION

We set out to test whether giving an LSTM an explicit per-sensor trend feature, computed with an ARIMA model, improves remaining useful life prediction on the NASA C-MAPSS benchmark. Under the fixed ARIMA order common in prior work, averaged over six seeds across all four sub-datasets, it does not: the hybrid model’s mean RMSE is at or above a plain LSTM’s everywhere, significantly so on FD004. Under a per-channel, AIC-selected order, tested at the same six-seed power, the picture changes: the selected-order hybrid significantly and consistently beats the fixed-order hybrid on two of four datasets, and on FD002 it significantly and consistently beats the plain LSTM baseline itself, the one result in this paper we consider a genuine, seed-robust architectural win rather than a null result or an artifact of insufficient replication. Getting to that answer took more checking than we originally planned on, an early, thinly-seeded version of the order-selection check looked unstable and nearly convinced us to report it as a third source of noise rather than a real effect, and only full six-seed power revealed it was not. We think the practical lesson generalizes past our own experiment: on a benchmark with 100–260 test engines, a single training run, or even a small handful of them, is not obviously enough to support a confident claim in either direction about whether an architectural change like this one helps, and we would encourage some caution toward papers in this space, including the two we cite favorably in Section II, that do not report how many seeds their headline numbers rest on. Future work includes a mechanistic account of why FD002 specifically benefits from a well-fit trend feature, testing whether an alternative fusion mechanism widens the benefit to FD001 and FD003, and validating the approach on proprietary industrial IoT sensor data from real manufacturing deployments.

ACKNOWLEDGMENT

Portions of the experimental code and manuscript drafting for this paper were produced with the assistance of a generative AI system, under the direct supervision and review of the author, who verified all reported experimental results, code, and claims. All experiments were executed on the publicly available NASA C-MAPSS dataset with no fabricated or synthetic result values.

REFERENCES

- [1] A. Saxena, K. Goebel, D. Simon, and N. Eklund, "Damage propagation modeling for aircraft engine run-to-failure simulation," in *Proc. Int. Conf. Prognostics and Health Management (PHM)*, 2008, pp. 1–9.
- [2] E. Ramasso and A. Saxena, "Performance benchmarking and analysis of prognostic methods for CMAPSS datasets," *Int. J. Prognostics and Health Management*, vol. 5, no. 2, pp. 1–15, 2014.
- [3] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.
- [4] X. Li, Q. Ding, and J.-Q. Sun, "Remaining useful life estimation in prognostics using deep convolution neural networks," *Reliability Engineering & System Safety*, vol. 172, pp. 1–11, 2018.
- [5] Muthukumar G. and J. Philip, "CNN-LSTM hybrid deep learning model for remaining useful life estimation," *arXiv:2412.15998*, Dec. 2024.
- [6] R. H. Chowdhury, N. Daiyan, F. Ahmed, M. R. Iqbal, and M. Sheikh, "Bi-cLSTM: residual-corrected bidirectional LSTM for aero-engine RUL estimation," *arXiv:2603.00745*, Mar. 2026.
- [7] Y. Su and X. Fang, "A multi-head attention fusion network for industrial prognostics under discrete operational conditions," *arXiv:2604.10248*, Apr. 2026.
- [8] S. Nie, Q. Liu, H. Ji, R. Hong, and S. Nie, "Integration of ARIMA and LSTM models for remaining useful life prediction of a water hydraulic high-speed on/off valve," *Applied Sciences*, vol. 12, no. 16, art. 8071, 2022.
- [9] Y. Liu, H. Zhang, and L. Wang, "ARIMA and attention-based CNN-LSTM hybrid neural network for battery life estimation," in *Proc. 2024 8th Int. Conf. Machine Learning and Soft Computing (ICMLSC)*, 2024, doi: 10.1145/3647750.3647759.
- [10] U. Khan, D. Cheng, F. Setti, F. Fummi, M. Cristani, and L. Capogrosso, "A comprehensive survey on deep learning-based predictive maintenance," *ACM Trans. Embedded Computing Systems*, vol. 25, no. 2, Apr. 2025.
- [11] K. Sharma, "Uncertainty-aware deep learning framework for remaining useful life prediction in turbofan engines with learned aleatoric uncertainty," *arXiv:2511.19124*, Nov. 2025.
- [12] P. Henderson, R. Islam, P. Bachman, J. Pineau, D. Precup, and D. Meger, "Deep reinforcement learning that matters," in *Proc. AAAI Conf. Artificial Intelligence*, 2018, pp. 3207–3214.
- [13] Z. C. Lipton and J. Steinhardt, "Troubling trends in machine learning scholarship," *arXiv:1807.03341*, Jul. 2018.
- [14] X. Bouthillier, C. Laurent, and P. Vincent, "Unreproducible research is reproducible," in *Proc. 36th Int. Conf. Machine Learning (ICML)*, 2019, pp. 725–734.
- [15] A. Goel, S. Galchar, and S. Kanu, "Remaining useful life estimation for turbofan engines: a comparative study of classical, CNN, and LSTM approaches," *arXiv:2604.27234*, Apr. 2026.